

EnsembleAssn.R

SeanOMalley1

Wed May 6 10:18:11 2015

```
# Ensemble Models in R
# Sean O'Malley

bank <- read.csv("/Users/SeanOMalley1/Desktop/Week\ 7\ ADM/bank-full.csv")

library(lattice)
library(ggplot2)
library(gplots)
```

```
##
## Attaching package: 'gplots'
##
## The following object is masked from 'package:stats':
##
##      lowess
```

```
library(mlbench)
library(plyr)
library(datasets)
library(graphics)
library(grDevices)
library(methods)
library(stats)
library(utils)
library(caret)
library(rpart)
library(randomForest)
```

```
## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(ROCR)
```

```
## Warning: package 'ROCR' was built under R version 3.1.3
```

```
library(gmodels)
library(doSNOW)
```

```
## Loading required package: foreach
## Loading required package: iterators
## Loading required package: snow
```

```
library(adabag)
library(rpart.plot)
```

```
# We are determining wheather someone is subscribing to a fixed term deposit to our bank. The classification model we are creating is going to determine the question, "What type of people subscribe to fixed term deposits?"
# Dependant Variable is "subscribed" variable
#The advantage of such a deposit is that the bank doesn't have to worry about the individual taking any money out for a fixed amount of time, thus providing more financial options and guaranteeing the money will be available to the bank for a longer period of time.

# EDA
str(bank) # classification random forest ensemble model where the dependant variable is y..category
```

```
## 'data.frame': 45211 obs. of 17 variables:
## $ age : int 58 44 33 47 33 35 28 42 58 43 ...
## $ job : Factor w/ 12 levels "admin.,""blue-collar",...: 5 10 3 2 12 5 5 3 6 10 ...
## $ marital : Factor w/ 3 levels "divorced","married",...: 2 3 2 2 3 2 3 1 2 3 ...
## $ education : Factor w/ 4 levels "primary","secondary",...: 3 2 2 4 4 3 3 3 1 2 ..
.
## $ default : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 2 1 1 ...
## $ balance : int 2143 29 2 1506 1 231 447 2 121 593 ...
## $ housing : Factor w/ 2 levels "no","yes": 2 2 2 2 1 2 2 2 2 2 ...
## $ loan : Factor w/ 2 levels "no","yes": 1 1 2 1 1 1 2 1 1 1 ...
## $ contact : Factor w/ 3 levels "cellular","telephone",...: 3 3 3 3 3 3 3 3 3 3 .
..
## $ day : int 5 5 5 5 5 5 5 5 5 5 ...
## $ month : Factor w/ 12 levels "apr","aug","dec",...: 9 9 9 9 9 9 9 9 9 9 ...
## $ duration : int 261 151 76 92 198 139 217 380 50 55 ...
## $ campaign : int 1 1 1 1 1 1 1 1 1 1 ...
## $ pdays : int -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 ...
## $ previous : int 0 0 0 0 0 0 0 0 0 0 ...
## $ outcome : Factor w/ 4 levels "failure","other",...: 4 4 4 4 4 4 4 4 4 4 ...
## $ subscribed: Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
```

```
names(bank)
```

```
## [1] "age"      "job"      "marital"  "education" "default"
## [6] "balance"  "housing"  "loan"     "contact"   "day"
## [11] "month"    "duration" "campaign" "pdays"    "previous"
## [16] "outcome"  "subscribed"
```

```
summary(bank) # many unknown variables
```

```
##          age                job                marital                education
## Min.      :18.00    blue-collar:9732    divorced: 5207    primary   : 6851
## 1st Qu.:33.00    management :9458    married :27214    secondary:23202
## Median :39.00    technician :7597    single  :12790    tertiary :13301
## Mean      :40.94    admin.      :5171                                unknown  : 1857
## 3rd Qu.:48.00    services    :4154
## Max.      :95.00    retired     :2264
##                (Other)    :6835
## default      balance      housing      loan                contact
## no :44396    Min.      : -8019    no :20081    no :37967    cellular :29285
## yes: 815    1st Qu.:      72    yes:25130    yes: 7244    telephone: 2906
##                Median :      448                                unknown  :13020
##                Mean      : 1362
##                3rd Qu.: 1428
##                Max.      :102127
##
##          day                month                duration                campaign
## Min.      : 1.00    may      :13766    Min.      : 0.0    Min.      : 1.000
## 1st Qu.: 8.00    jul      : 6895    1st Qu.: 103.0    1st Qu.: 1.000
## Median :16.00    aug      : 6247    Median : 180.0    Median : 2.000
## Mean      :15.81    jun      : 5341    Mean      : 258.2    Mean      : 2.764
## 3rd Qu.:21.00    nov      : 3970    3rd Qu.: 319.0    3rd Qu.: 3.000
## Max.      :31.00    apr      : 2932    Max.      :4918.0    Max.      :63.000
##                (Other): 6060
##          pdays      previous      outcome      subscribed
## Min.      : -1.0    Min.      : 0.0000    failure: 4901    no :39922
## 1st Qu.: -1.0    1st Qu.: 0.0000    other : 1840    yes: 5289
## Median : -1.0    Median : 0.0000    success: 1511
## Mean      : 40.2    Mean      : 0.5803    unknown:36959
## 3rd Qu.: -1.0    3rd Qu.: 0.0000
## Max.      :871.0    Max.      :275.0000
##
```

```
#####
```

```
# Model 1: Dec Tree # All Variables
```

```
  # Train and Test/Holdout Examples
```

```
set.seed(99)
```

```
bank_rand1 <- bank[order(runif(45211)),] #45,211 total observations
```

```
bank_rand1 <- sample(1:45211, 31648) #want 31,648 obs for training and 13,563 obs in  
test. that is about a 70/30
```

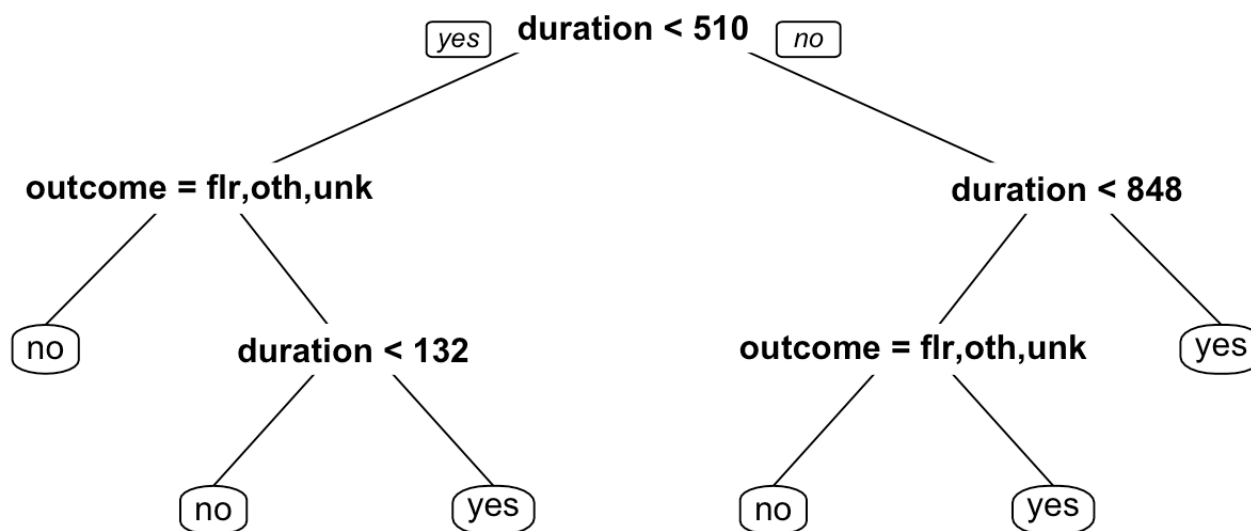
```
bank_train1 <- bank[ bank_rand1,]
```

```
bank_test1  <- bank[-bank_rand1,]
```

```
# Decision Tree: # (rpart)
```

```
bank_tree1 <- rpart(subscribed~.,data=bank_train1)
```

```
rpart.plot(bank_tree1)
```



```
print(bank_tree1)
```

```
## n= 31648
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 31648 3710 no (0.88277300 0.11722700)
##    2) duration< 510.5 28071 2134 no (0.92397848 0.07602152)
##      4) outcome=failure,other,unknown 27171 1581 no (0.94181296 0.05818704) *
##      5) outcome=success 900 347 yes (0.38555556 0.61444444)
##        10) duration< 132.5 178 43 no (0.75842697 0.24157303) *
##        11) duration>=132.5 722 212 yes (0.29362881 0.70637119) *
##    3) duration>=510.5 3577 1576 no (0.55940732 0.44059268)
##      6) duration< 847.5 2401 880 no (0.63348605 0.36651395)
##        12) outcome=failure,other,unknown 2281 781 no (0.65760631 0.34239369) *
##        13) outcome=success 120 21 yes (0.17500000 0.82500000) *
##    7) duration>=847.5 1176 480 yes (0.40816327 0.59183673) *
```

summary(bank_tree1) # shows duration and outcome as only important factors, could mean parameter adjustment needed

```
## Call:
## rpart(formula = subscribed ~ ., data = bank_train1)
##      n= 31648
##
##              CP nsplit rel error   xerror   xstd
## 1 0.03791554      0 1.0000000 1.0000000 0.01542544
## 2 0.02479784      3 0.8862534 0.8900270 0.01465843
## 3 0.02102426      4 0.8614555 0.8638814 0.01446621
## 4 0.01000000      5 0.8404313 0.8455526 0.01432902
##
## Variable importance
## duration  outcome
##      62      37
##
## Node number 1: 31648 observations,      complexity param=0.03791554
## predicted class=no expected loss=0.117227 P(node) =1
## class counts: 27938 3710
## probabilities: 0.883 0.117
## left son=2 (28071 obs) right son=3 (3577 obs)
## Primary splits:
##      duration < 510.5 to the left,      improve=843.3836, (0 missing)
##      outcome splits as LLRL,      improve=622.3676, (0 missing)
##      month splits as LLRLLLLRLLRR, improve=335.0443, (0 missing)
##      pdays < 8.5 to the left,      improve=190.5850, (0 missing)
##      previous < 0.5 to the left,      improve=187.1104, (0 missing)
```

```

##
## Node number 2: 28071 observations,      complexity param=0.03791554
## predicted class=no expected loss=0.07602152 P(node) =0.8869755
## class counts: 25937 2134
## probabilities: 0.924 0.076
## left son=4 (27171 obs) right son=5 (900 obs)
## Primary splits:
## outcome splits as LLRL, improve=539.1031, (0 missing)
## month splits as LLRLLLLRLRLRR, improve=314.3254, (0 missing)
## pdays < 16 to the left, improve=173.0977, (0 missing)
## previous < 0.5 to the left, improve=169.8065, (0 missing)
## duration < 206.5 to the left, improve=146.7630, (0 missing)
##
## Node number 3: 3577 observations,      complexity param=0.03791554
## predicted class=no expected loss=0.4405927 P(node) =0.1130245
## class counts: 2001 1576
## probabilities: 0.559 0.441
## left son=6 (2401 obs) right son=7 (1176 obs)
## Primary splits:
## duration < 847.5 to the left, improve=80.15318, (0 missing)
## outcome splits as LLRL, improve=51.35249, (0 missing)
## contact splits as RRL, improve=46.33695, (0 missing)
## month splits as LLRLLLLRLRLRR, improve=35.23122, (0 missing)
## pdays < 8.5 to the left, improve=22.24908, (0 missing)
## Surrogate splits:
## campaign < 22.5 to the left, agree=0.672, adj=0.003, (0 split)
## previous < 17.5 to the left, agree=0.672, adj=0.003, (0 split)
## age < 87.5 to the left, agree=0.672, adj=0.001, (0 split)
## balance < -1207 to the right, agree=0.672, adj=0.001, (0 split)
## pdays < 392.5 to the left, agree=0.672, adj=0.001, (0 split)
##
## Node number 4: 27171 observations
## predicted class=no expected loss=0.05818704 P(node) =0.8585377
## class counts: 25590 1581
## probabilities: 0.942 0.058
##
## Node number 5: 900 observations,      complexity param=0.02479784
## predicted class=yes expected loss=0.3855556 P(node) =0.02843782
## class counts: 347 553
## probabilities: 0.386 0.614
## left son=10 (178 obs) right son=11 (722 obs)
## Primary splits:
## duration < 132.5 to the left, improve=61.698340, (0 missing)
## housing splits as RL, improve=11.996250, (0 missing)
## month splits as RRRRRRRRLRLRR, improve= 8.053067, (0 missing)
## pdays < 51.5 to the left, improve= 7.023146, (0 missing)
## campaign < 3.5 to the right, improve= 5.919221, (0 missing)

```

```
## Surrogate splits:
##   contact splits as RRL, agree=0.810, adj=0.039, (0 split)
##   default splits as RL, agree=0.803, adj=0.006, (0 split)
##
## Node number 6: 2401 observations,   complexity param=0.02102426
##   predicted class=no   expected loss=0.366514   P(node) =0.07586577
##   class counts:  1521   880
##   probabilities: 0.633 0.367
##   left son=12 (2281 obs) right son=13 (120 obs)
##   Primary splits:
##   outcome splits as LLRL,           improve=53.10438, (0 missing)
##   contact splits as RRL,           improve=35.94764, (0 missing)
##   month splits as LRLLLLRLLRR, improve=32.56742, (0 missing)
##   pdays < 8.5 to the left, improve=27.50869, (0 missing)
##   previous < 0.5 to the left, improve=27.39978, (0 missing)
##
## Node number 7: 1176 observations
##   predicted class=yes expected loss=0.4081633   P(node) =0.03715875
##   class counts:   480   696
##   probabilities: 0.408 0.592
##
## Node number 10: 178 observations
##   predicted class=no   expected loss=0.241573   P(node) =0.005624368
##   class counts:   135   43
##   probabilities: 0.758 0.242
##
## Node number 11: 722 observations
##   predicted class=yes expected loss=0.2936288   P(node) =0.02281345
##   class counts:   212   510
##   probabilities: 0.294 0.706
##
## Node number 12: 2281 observations
##   predicted class=no   expected loss=0.3423937   P(node) =0.07207406
##   class counts:  1500   781
##   probabilities: 0.658 0.342
##
## Node number 13: 120 observations
##   predicted class=yes expected loss=0.175   P(node) =0.003791709
##   class counts:    21   99
##   probabilities: 0.175 0.825
```

```
#confusion matrix for rpart
bank_tree1_actual <- bank_test1$subscribed #created to test the "test" data/
bank_tree1_pred <- predict(bank_tree1, bank_test1, type="class")
bank_tree1_results <- confusionMatrix(bank_tree1_pred, bank_tree1_actual) #the model
vs the actual holdout data.
print(bank_tree1_results)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    no   yes
##           no 11668 1033
##           yes  316  546
##
##           Accuracy : 0.9005
##           95% CI : (0.8954, 0.9055)
##           No Information Rate : 0.8836
##           P-Value [Acc > NIR] : 1.705e-10
##
##           Kappa : 0.3978
##           McNemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.9736
##           Specificity : 0.3458
##           Pos Pred Value : 0.9187
##           Neg Pred Value : 0.6334
##           Prevalence : 0.8836
##           Detection Rate : 0.8603
##           Detection Prevalence : 0.9364
##           Balanced Accuracy : 0.6597
##
##           'Positive' Class : no
##
```



```
# We have 90% accuracy at predicting whether someone will subscribe to fixed term deposits
# Greedy algorithm is only using 2/17 variables available to make a prediction, need further research
```

```
#####
```

```
# Model 2 : Random Forest
set.seed(99)
bank_rfl <- randomForest(subscribed ~.,data=bank_train1, ntree=25,na.action = na.omit,
, importance=TRUE)
# can put ~.because he got rid of id #, see line 6
print(bank_rfl) #shows OOB of model and confusion matrix
```

```
##
## Call:
## randomForest(formula = subscribed ~ ., data = bank_train1, ntree = 25,      impor
tance = TRUE, na.action = na.omit)
##              Type of random forest: classification
##              Number of trees: 25
## No. of variables tried at each split: 4
##
##              OOB estimate of  error rate: 9.91%
## Confusion matrix:
##              no  yes class.error
## no   26710 1227  0.04392025
## yes   1910 1800  0.51482480
```

```
importance(bank_rfl) #shows the importance of each variable
```

##		no	yes	MeanDecreaseAccuracy	MeanDecreaseGini
##	age	10.063115	4.7133183	12.075863	571.611929
##	job	7.134809	0.2691294	7.389245	433.964363
##	marital	1.245914	2.9150187	2.539107	118.530072
##	education	5.753903	0.7716123	5.475338	152.246258
##	default	1.071415	1.9177811	2.085381	9.729627
##	balance	3.892177	2.3841880	4.489242	617.421018
##	housing	10.222948	6.3030120	10.504654	138.627167
##	loan	1.037465	3.2801608	2.851221	47.517952
##	contact	13.708937	1.5313865	14.431085	120.340680
##	day	16.939956	2.5380681	16.578163	521.310875
##	month	27.221931	7.7196258	30.402505	731.948972
##	duration	27.755439	60.0978345	50.333508	1828.074949
##	campaign	5.165356	2.6466831	5.879911	235.754048
##	pdays	4.807660	5.1166029	5.384774	270.696341
##	previous	5.266209	3.4538595	5.521918	152.959713
##	outcome	7.859392	3.4875573	11.185021	415.124405

No. of variables tried at each split: 4 # at each branch it picks 4/17 variables, then picks best one for tree.

then after the trees are ensemble it takes the best of 500 trees that are picky

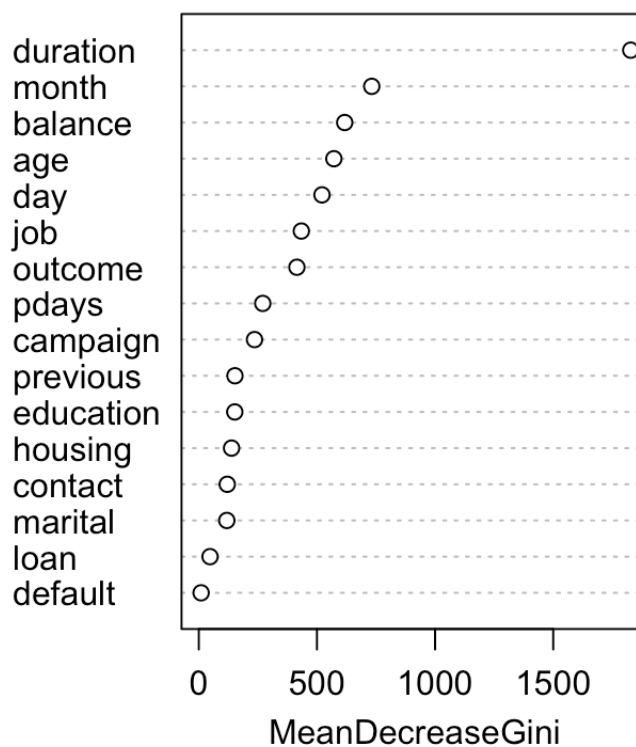
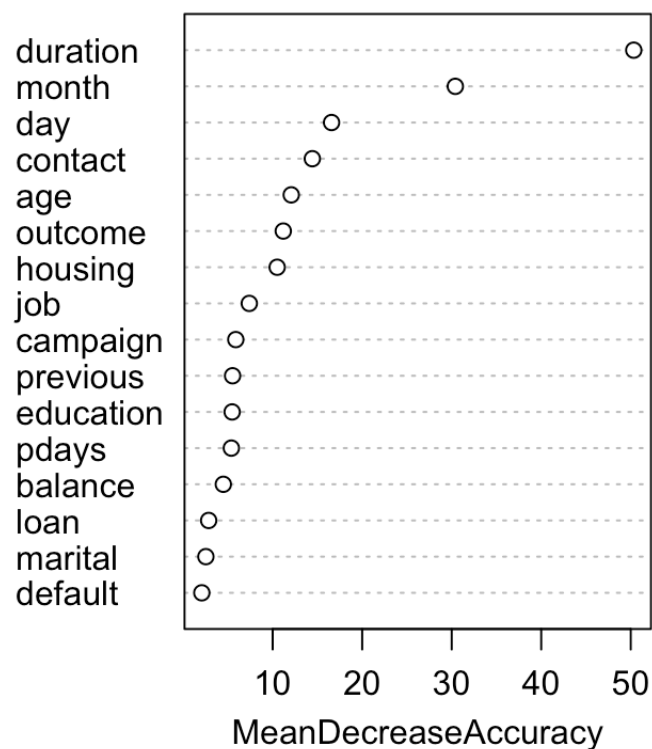
OOB estimate of error rate: tells us that the average error rate of all trees was 9.31% aka 90.69% accuracy

confusion matrix inaccuracy shows this is all run off of training data, yet to do test data

unbalanced dataset with majority being true negatives

varImpPlot(bank_rfl) #plots the importance of each variable

bank_rf1



```
# unlike the decision tree, default, loan and balance affect the accuracy of the dv t
he most
# Gini measures entropy, accuracy is most important, further left = more important
# tells us most important features to tell manager this way, doesn't give split point
s bc theyre synthetic
```

```
# Running test through test data
bank_rf1_actual <- bank_test1$subscribed
bank_rf1_pred <- predict(bank_rf1, bank_test1, type="response")
bank_rf1_results <- confusionMatrix(bank_rf1_pred, bank_rf1_actual)
print(bank_rf1_results)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    no    yes
##           no 11509   795
##           yes  475   784
##
##           Accuracy : 0.9064
##           95% CI : (0.9013, 0.9112)
##           No Information Rate : 0.8836
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.501
##           McNemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.9604
##           Specificity : 0.4965
##           Pos Pred Value : 0.9354
##           Neg Pred Value : 0.6227
##           Prevalence : 0.8836
##           Detection Rate : 0.8486
##           Detection Prevalence : 0.9072
##           Balanced Accuracy : 0.7284
##
##           'Positive' Class : no
##
```

```
# accuracy improves marginally, but kappa and sensitivity show the greatest improvements
# making this model better than the tree model
CrossTable(bank_rfl_pred, bank_rfl_actual)
```

```
##
##
##      Cell Contents
## |-----|
## |                      N |
## | Chi-square contribution |
## |      N / Row Total    |
## |      N / Col Total    |
## |      N / Table Total  |
## |-----|
##
##
## Total Observations in Table:  13563
##
##
##      | bank_rfl_actual
## bank_rfl_pred |      no      |      yes      | Row Total |
## -----|-----|-----|-----|
##           no |    11509     |      795      |    12304  |
##           |    37.374    |    283.654    |           |
##           |     0.935    |     0.065     |    0.907  |
##           |     0.960    |     0.503     |           |
##           |     0.849    |     0.059     |           |
## -----|-----|-----|-----|
##           yes |      475     |      784      |    1259   |
##           |    365.250   |    2772.105   |           |
##           |     0.377   |     0.623     |    0.093  |
##           |     0.040   |     0.497     |           |
##           |     0.035   |     0.058     |           |
## -----|-----|-----|-----|
## Column Total |    11984     |    1579       |    13563  |
##           |     0.884     |     0.116     |           |
## -----|-----|-----|-----|
##
##
```

```

# true positives and negatives show great increase in comparison to previous model

# K Fold Cross Validation of Model 2
registerDoSNOW(makeCluster(2, type = "SOCK")) #sets to use the 2 cores in my laptop
ctrl <- trainControl(method = "repeatedcv",
                     number = 5, repeats = 5) #sets crossvalidation run kfold against
itself
# creates kfold of 10 # this syntax essentially just sets tuning 10^10
# each fold and each cross validation is parallelizable

set.seed(99)
bank_rfl_cv <- train(subscribed ~., data=bank_train1, method = "rf",
                    metric = "Kappa", trControl = ctrl) #method is randomforest.
#Adds in the ctrl as control of the line above
# takes 500 trees and randomly tries multiple tries at trees

print(bank_rfl_cv)

```

```

## Random Forest
##
## 31648 samples
##    16 predictor
##    2 classes: 'no', 'yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
##
## Summary of sample sizes: 25318, 25319, 25318, 25319, 25318, 25319, ...
##
## Resampling results across tuning parameters:
##
##  mtry  Accuracy   Kappa     Accuracy SD   Kappa SD
##    2    0.8924039 0.1682657 0.001006159   0.01345966
##   22    0.9048786 0.4805970 0.002650910   0.01814799
##   42    0.9042593 0.4819962 0.003013329   0.01994296
##
## Kappa was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 42.

```

```

# mtry is the approximate square of how many splits
# showed us random splits that gave us the highest kappas, which is what we asked for
# grid search: goes through a range of # and find what works best
#use all your other RF diagnostics and plots here
    # 5 splits (22mtry) is the optimal amount of kappa at 0.483, with 0.905 accuracy

# Auto-tune random forest
bank_rfl_grid <- expand.grid(.mtry = c(4,9,25))
#mtry is how many variables it randomly tries at each node (sqrt of mtry is # of splits)

set.seed(99)
bank_ptm <- proc.time() #starts to time
bank_rfl_cv2 <- train(subscribed ~.,data=bank_train1, method = "rf", ntree=500,
                      metric = "Kappa", trControl = ctrl, tuneGrid = bank_rfl_grid)
#also control mtry
print(bank_rfl_cv2)

```

```

## Random Forest
##
## 31648 samples
##    16 predictor
##    2 classes: 'no', 'yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 5 times)
##
## Summary of sample sizes: 25318, 25319, 25318, 25319, 25318, 25319, ...
##
## Resampling results across tuning parameters:
##
##   mtry  Accuracy   Kappa      Accuracy SD   Kappa SD
##    4    0.9025847  0.3592532  0.001548192   0.01712334
##    9    0.9057633  0.4679619  0.002682394   0.01870807
##   25    0.9048533  0.4822413  0.002975948   0.02019670
##
## Kappa was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 25.

```


```
# after adding ctrl and tune
# mtry is how many variables it randomly tries at each node (sqrt of mtry is # of splits)
# if you want the highest kappa, use the corresponding mtry to find optimal number of splits
# could go back and do a grid search to find best kmeans

summary(bank_rf1_cv2)
```

##	Length	Class	Mode
## call	5	-none-	call
## type	1	-none-	character
## predicted	31648	factor	numeric
## err.rate	1500	-none-	numeric
## confusion	6	-none-	numeric
## votes	63296	matrix	numeric
## oob.times	31648	-none-	numeric
## classes	2	-none-	character
## importance	42	-none-	numeric
## importanceSD	0	-none-	NULL
## localImportance	0	-none-	NULL
## proximity	0	-none-	NULL
## ntree	1	-none-	numeric
## mtry	1	-none-	numeric
## forest	14	-none-	list
## y	31648	factor	numeric
## test	0	-none-	NULL
## inbag	0	-none-	NULL
## xNames	42	-none-	character
## problemType	1	-none-	character
## tuneValue	1	data.frame	list
## obsLevels	2	-none-	character

```
# ROC Curve for Decision Tree
bank_tree1_pred_prob1 <- predict(bank_tree1, type="prob", bank_test1) # same as above f
or predict, but add "prob".
bank_pred2 <- prediction(bank_tree1_pred_prob1[,2], bank_test1$subscribed)
bank_perf2 <- performance(bank_pred2, "tpr", "fpr") #true pos and false pos
plot(bank_perf2, main="ROC Curve for Decision Tree", col=2, lwd=2)
abline(a=0, b=1, lwd=2, lty=2, col="gray")
```


ROC Curve for Decision Tree



```
#area under the curve
bank_tree1_auc <- performance(bank_pred2, measure = "auc") #run an area under the curve
str(bank_tree1_auc) #see different values in the auc object
```

```
## Formal class 'performance' [package "ROCR"] with 6 slots
##   ..@ x.name      : chr "None"
##   ..@ y.name      : chr "Area under the ROC curve"
##   ..@ alpha.name   : chr "none"
##   ..@ x.values     : list()
##   ..@ y.values     : List of 1
##   .. ..$ : num 0.749
##   ..@ alpha.values: list()
```

```
as.numeric(bank_tree1_auc@y.values) #shows the AUC percentage
```

```
## [1] 0.7485122
```

```
# 74% above
```

```
# ROC curve for RF1
```

```
bank_rf1_pred_prob1 <- predict(bank_rf1, type="prob", bank_test1) # same as above for p
predict, but add "prob".
```


```
bank_pred1 <- prediction(bank_rf1_pred_prob1[,2], bank_test1$subscribed)
```

```
bank_perfl <- performance(bank_pred1, "tpr", "fpr") #true pos and false pos
```

```
plot(bank_perfl, main="ROC Curve for Random Forest (without tuning)", col=2, lwd=2)
```

```
abline(a=0, b=1, lwd=2, lty=2, col="gray")
```

ROC Curve for Random Forest (without tuning)



```
#area under the curve
```

```
bank_rf1_auc <- performance(bank_pred1, measure = "auc") #run an area under the curve
str(bank_rf1_auc) #see different values in the auc object
```

```
## Formal class 'performance' [package "ROCR"] with 6 slots
## ..@ x.name      : chr "None"
## ..@ y.name      : chr "Area under the ROC curve"
## ..@ alpha.name  : chr "none"
## ..@ x.values    : list()
## ..@ y.values    :List of 1
## .. ..$ : num 0.92
## ..@ alpha.values: list()
```

```
as.numeric(bank_rf1_auc@y.values) #shows the AUC percentage
```

```
## [1] 0.9202552
```

```
# 93% above
```

```
# ROC curve for RF2
```

```
bank_rf2_pred_prob1 <- predict(bank_rf1_cv, type="prob", bank_test1) # same as above for predict, but add "prob".
```


```
bank_pred3 <- prediction(bank_rf2_pred_prob1[,2], bank_test1$subscribed)
```

```
bank_perf3 <- performance(bank_pred3, "tpr", "fpr") #true pos and false pos
```

```
plot(bank_perf3, main="ROC Curve for Random Forest 2 (k=10)", col=2, lwd=2)
```

```
abline(a=0, b=1, lwd=2, lty=2, col="gray")
```

ROC Curve for Random Forest 2 (k=10)



#area under the curve

```
bank_rf2_auc <- performance(bank_pred3, measure = "auc") #run an area under the curve
str(bank_rf2_auc) #see different values in the auc object
```

```
## Formal class 'performance' [package "ROCR"] with 6 slots
##   ..@ x.name      : chr "None"
##   ..@ y.name      : chr "Area under the ROC curve"
##   ..@ alpha.name   : chr "none"
##   ..@ x.values     : list()
##   ..@ y.values     : List of 1
##   .. ..$ : num 0.924
##   ..@ alpha.values: list()
```

```
as.numeric(bank_rf2_auc@y.values) #shows the AUC percentage
```

```
## [1] 0.9239141
```

```
# 92.7% above (computationally expensive)
```

```
# ROC curve for RF3
```

```
bank_rf3_pred_prob1 <- predict(bank_rf1_cv2, type="prob", bank_test1) # same as above f
or predict, but add "prob".
```


```
bank_pred4 <- prediction(bank_rf3_pred_prob1[,2], bank_test1$subscribed)
```

```
bank_perf4 <- performance(bank_pred4, "tpr", "fpr") #true pos and false pos
```

```
plot(bank_perf4, main="ROC Curve for Random Forest 3 (auto-tuned)(k=10)", col=2, lwd=2)
```

```
abline(a=0, b=1, lwd=2, lty=2, col="gray")
```

ROC Curve for Random Forest 3 (auto-tuned)(k=10)



```
#area under the curve
```

```
bank_rf3_auc <- performance(bank_pred4, measure = "auc") #run an area under the curve
str(bank_rf3_auc) #see different values in the auc object
```

```
## Formal class 'performance' [package "ROCR"] with 6 slots
##   ..@ x.name      : chr "None"
##   ..@ y.name      : chr "Area under the ROC curve"
##   ..@ alpha.name   : chr "none"
##   ..@ x.values     : list()
##   ..@ y.values     :List of 1
##   .. ..$ : num 0.927
##   ..@ alpha.values: list()
```

```
as.numeric(bank_rf3_auc@y.values) #shows the AUC percentage
```

```
## [1] 0.9271857
```

```
# 92.6% above (computationally expensive)
```